

# AdaBoost: From Weak Learners to Strong Classifiers

## Minimizing Exponential Error via Sequential Learning

Nikhil, Jannen

January 21, 2026

# The Power of the Committee

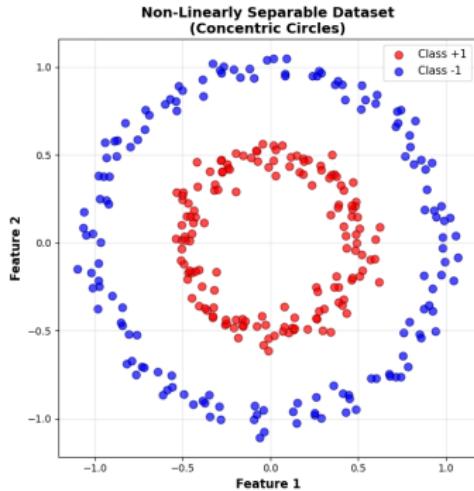


Figure: Linear Separability Challenge

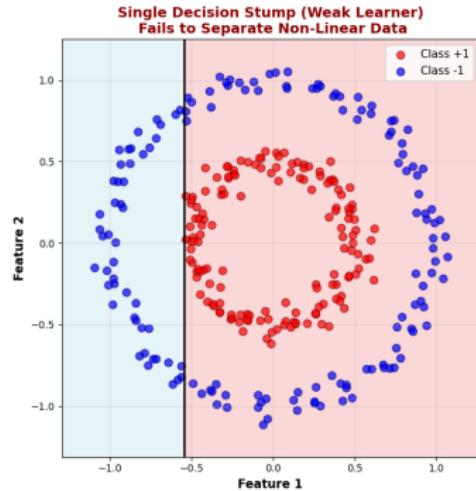


Figure: Decision Stump fails

- **The Problem:** Complex data is rarely linearly separable.
- **The Naive Solution:** Build a massive, complex model.
- **The AdaBoost Solution:** Combine simple “Decision boundaries.”

$$H(x) = \text{sign} \left( \sum \alpha_k h_k(x) \right)$$

# Why Not Squared Error?

- **Regression:** Minimizes Least Squared Error (LSE).
- **Classification:** LSE is unreliable.
  - Penalizes “too correct” predictions (Overfitting).
  - Sensitive to outliers in the wrong way.
- **AdaBoost:** Minimizes Exponential Error.

$$E = \sum e^{-t_n f(x_n)}$$

Focuses heavily on misclassified points ( $t \neq f(x)$ ).

# Sequential Minimization

$$E = \sum_{n=1}^N \exp(-t_n [H_{k-1}(x_n) + \alpha_k h_k(x_n)])$$

- Cannot optimize all  $\alpha$  and  $h$  at once.
- **Greedy Approach:** Freeze past, optimize current step.
- **The Weight Trick:**

$$w_n^{(k)} = \exp(-t_n H_{k-1}(x_n))$$

- Previous errors become current weights!
- Misclassified points get higher weights.

# Finding the Optimal Weight ( $\alpha_k$ )

- **Goal:** Find  $\alpha_k$  to minimize  $E$ .
- Taking the derivative:  $\frac{\partial E}{\partial \alpha_k} = 0$
- **Result:**

$$\alpha_k = \frac{1}{2} \ln \left( \frac{1 - \epsilon_k}{\epsilon_k} \right)$$

- Low Error ( $\epsilon_k \rightarrow 0$ )  $\Rightarrow$  High Alpha
- Random Guess ( $\epsilon_k = 0.5$ )  $\Rightarrow$  Zero Alpha
- Worse than random ( $\epsilon_k > 0.5$ )  $\Rightarrow$  Negative Alpha

# Updating Sample Weights

- After computing  $\alpha_k$  and training  $h_k$ , update weights:

$$w_n^{(k+1)} = w_n^{(k)} \cdot \exp(-\alpha_k \cdot t_n \cdot h_k(x_n))$$

- If correct:**  $t_n \cdot h_k(x_n) = 1 \Rightarrow$  weight decreases
- If wrong:**  $t_n \cdot h_k(x_n) = -1 \Rightarrow$  weight increases
- Normalize:**  $w_n^{(k+1)} \leftarrow \frac{w_n^{(k+1)}}{\sum_i w_i^{(k+1)}}$

Hard examples get harder weights: forces next learner to focus.

# The AdaBoost Algorithm

① **Initialize:**  $w_n^{(1)} = \frac{1}{N}$

② **For**  $k = 1, 2, \dots, K$ :

    ① Train weak learner  $h_k$  on weighted data

    ② Calculate error:  $\epsilon_k = \frac{\sum_n w_n^{(k)} \cdot I(h_k(x_n) \neq t_n)}{\sum_i w_i^{(k)}}$

    ③ Compute weight:  $\alpha_k = \frac{1}{2} \ln \left( \frac{1-\epsilon_k}{\epsilon_k} \right)$

    ④ Update:  $w_n^{(k+1)} = w_n^{(k)} \cdot \exp(-\alpha_k \cdot t_n \cdot h_k(x_n))$

    ⑤ Normalize:  $w_n^{(k+1)} \leftarrow \frac{w_n^{(k+1)}}{\sum_i w_i^{(k+1)}}$

③ **Output:**  $H(x) = \text{sign} \left( \sum_{k=1}^K \alpha_k h_k(x) \right)$

# The Dataset: Breast Cancer Wisconsin (Diagnostic)

## What are we classifying?

- **Source:** FNA biopsy images
- **Samples:** 569 patients
- **Task:** Binary classification
- Malignant (M): ~ 37%
- Benign (B): ~ 63%

## Why it matters:

- Early cancer detection
- Non-invasive procedure
- Real-world ML

## Feature Engineering:

- 10 cell nucleus characteristics
  - 3 measurements each
  - ⇒ 30 total features
- 1 Radius
  - 2 Texture
  - 3 Perimeter
  - 4 Area
  - 5 Smoothness
  - 6 Compactness
  - 7 Concavity
  - 8 Concave Points
  - 9 Symmetry
  - 10 Fractal Dimension

# Visualization Strategy: 30D to 2D Projection

## The Challenge:

- 30 features  $\Rightarrow$  impossible to visualize
- Need 2D representation

## Solution: PCA

- Projects 30D to 2D
- Preserves maximum variance
- Classes remain **non-linearly separable**

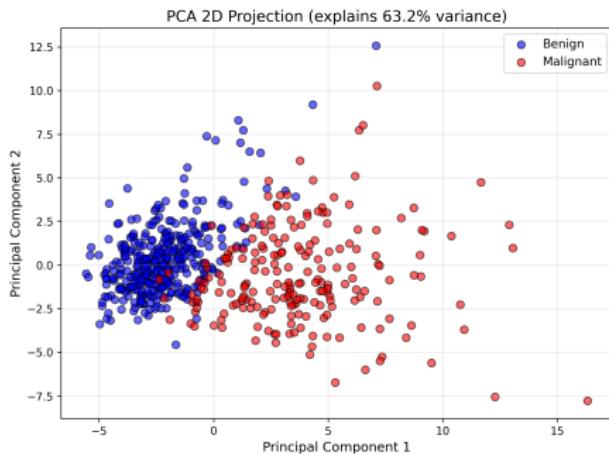


Figure: 2D projection of 30D breast cancer features

# Live Demo: AdaBoost on Breast Cancer (2D PCA)

## What to Observe:

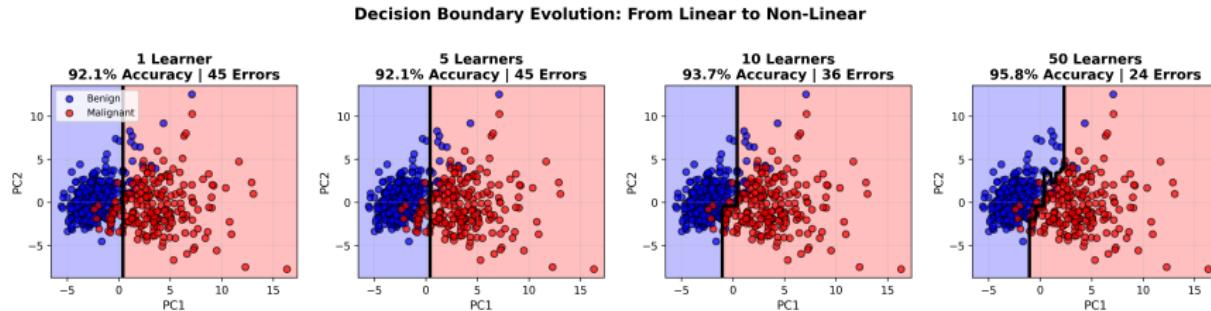
- Decision boundary starts simple (linear stump)
- Boundary becomes increasingly complex
- Classification accuracy improves with iterations
- Misclassified points get higher weights
- Final ensemble combines weak learners effectively

## Expected Behavior:

- ➊ **Iteration 1:** Single decision stump
- ➋ **Iterations 2-5:** Boundary refines, adapts to errors
- ➌ **Iterations 6-10:** Complex non-linear boundary
- ➍ **Final:** Accurate classification on 2D PCA space

*The algorithm iteratively focuses on hard to classify samples*

# Why AdaBoost Works: The Exponential Error



**Figure:** Decision boundary evolution: 1, 5, 10, 50 weak learners. Blue regions = Benign class, Red regions = Malignant class. Black line = decision boundary.

- **1 Learner:** Single stump creates simple linear split
  - Makes one decision: "if feature > threshold"
  - Cannot capture non-linear patterns
  - Accuracy:  $\sim 92\%$  — 45 errors
- **5-50 Learners:** Ensemble refines boundary iteratively
  - Error weights grow exponentially:  $w \leftarrow w \cdot e^{-\alpha y h(x)}$
  - Each stump targets previous mistakes

# Summary

- AdaBoost minimizes Exponential Error
- Turns complex problem into simple sequence (greedy)
- Weight update: Focuses on misclassified examples
- Key Takeaway: Smart combination of simple models
- Rigorous: Every formula from exponential loss

**Github:** [https://github.com/Jannen06/Adaboost\\_ML\\_nonlineair](https://github.com/Jannen06/Adaboost_ML_nonlineair)

**AdaBoost: Mathematically principled and empirically powerful**

# References & Tools

- **Literature:**

- Bishop, C. M. (2006). *Pattern Recognition* (Ch. 14).
- Friedman, Hastie, Tibshirani (2000). *Additive Logistic Regression*.
- Viola, Jones (2004). *Robust Real-Time Face Detection*.

- **Dataset:**

- Wolberg et al. (1995). *Breast Cancer Wisconsin*. UCI ML. DOI: 10.24432/C5DW2B
- Dataset Link

- **Tools:**

- scikit-learn, ucimlrepo, matplotlib